

Table of Contents

 Learn

 Colab

 Notebook

 GitHub

[Learn the Basics](#) || [Quickstart](#) || [Tensors](#) || [Datasets & DataLoaders](#) || [Transforms](#) || **[Build Model](#)** || [Autograd](#) || [Optimization](#) || [Save & Load Model](#)

BUILD THE NEURAL NETWORK

Neural networks comprise of layers/modules that perform operations on data. The `torch.nn` namespace provides all the building blocks you need to build your own neural network. Every module in PyTorch subclasses the `nn.Module`. A neural network is a module itself that consists of other modules (layers). This nested structure allows for building and managing complex architectures easily.

In the following sections, we'll build a neural network to classify images in the FashionMNIST dataset.

```
import os
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
```

Get Device for Training

We want to be able to train our model on a hardware accelerator like the GPU, if it is available. Let's check to see if `torch.cuda` is available, else we continue to use the CPU.

```
device = "cuda" if torch.cuda.is_available() else "cpu"
print(f"Using {device} device")
```

Out:

Using cuda device

Define the Class

We define our neural network by subclassing `nn.Module`, and initialize the neural network layers in `__init__`. Every `nn.Module` subclass implements the operations on input data in the `forward` method.

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```

We create an instance of `NeuralNetwork`, and move it to the `device`, and print its structure.

```
model = NeuralNetwork().to(device)
print(model)
```

Out:

```
NeuralNetwork(  
    (flatten): Flatten(start_dim=1, end_dim=-1)  
    (linear_relu_stack): Sequential(  
      (0): Linear(in_features=784, out_features=512, bias=True)  
      (1): ReLU()  
      (2): Linear(in_features=512, out_features=512, bias=True)  
      (3): ReLU()  
      (4): Linear(in_features=512, out_features=10, bias=True)  
    )  
  )  
)
```

To use the model, we pass it the input data. This executes the model's `forward`, along with some **background operations**. Do not call `model.forward()` directly!

Calling the model on the input returns a 10-dimensional tensor with raw predicted values for each class. We get the prediction probabilities by passing it through an instance of the `nn.Softmax` module.

```
X = torch.rand(1, 28, 28, device=device)  
logits = model(X)  
pred_probab = nn.Softmax(dim=1)(logits)  
y_pred = pred_probab.argmax(1)  
print(f"Predicted class: {y_pred}")
```

Out:

```
Predicted class: tensor([7], device='cuda:0')
```

Model Layers

Let's break down the layers in the FashionMNIST model. To illustrate it, we will take a sample minibatch of 3 images of size 28x28 and see what happens to it as we pass it through the network.

```
input_image = torch.rand(3,28,28)  
print(input_image.size())
```

Out:

```
torch.Size([3, 28, 28])
```

nn.Flatten

We initialize the **nn.Flatten** layer to convert each 2D 28x28 image into a contiguous array of 784 pixel values (the minibatch dimension (at dim=0) is maintained).

```
flatten = nn.Flatten()  
flat_image = flatten(input_image)  
print(flat_image.size())
```

Out:

```
torch.Size([3, 784])
```

nn.Linear

The **linear layer** is a module that applies a linear transformation on the input using its stored weights and biases.

```
layer1 = nn.Linear(in_features=28*28, out_features=20)  
hidden1 = layer1(flat_image)  
print(hidden1.size())
```

Out:

```
torch.Size([3, 20])
```

nn.ReLU

Non-linear activations are what create the complex mappings between the model’s inputs and outputs. They are applied after linear transformations to introduce *nonlinearity*, helping neural networks learn a wide variety of phenomena.

In this model, we use **nn.ReLU** between our linear layers, but there’s other activations to introduce non-linearity in your model.

```
print(f"Before ReLU: {hidden1}\n\n")
hidden1 = nn.ReLU()(hidden1)
print(f"After ReLU: {hidden1}")
```

Out:

```
Before ReLU: tensor([[ -0.2473,  0.2972, -0.0437, -0.3515,  0.1597, -0.2566,  0.3498, -0.0339,
                    -0.1560, -0.3464, -0.0351,  0.5189,  0.2862, -0.1267,  0.1509, -0.3907,
                    -0.4598, -0.1547, -0.0264, -0.5889],
                    [-0.3076, -0.3013, -0.3605, -0.4714,  0.2394,  0.0326,  0.4283,  0.2156,
                    -0.1660, -0.5472, -0.1775,  0.6659,  0.1813, -0.4881,  0.1295, -0.4282,
                    -0.4571, -0.3703,  0.1829, -0.4710],
                    [-0.3813,  0.1029, -0.1572, -0.6220, -0.0048,  0.3352, -0.0507, -0.1262,
                    -0.0825, -0.6460,  0.0912,  0.5666, -0.0549, -0.3259,  0.0430, -0.1132,
                    -0.5483,  0.1878, -0.1335, -0.5314]], grad_fn=<AddmmBackward0>)

After ReLU: tensor([[0.0000, 0.2972, 0.0000, 0.0000, 0.1597, 0.0000, 0.3498, 0.0000, 0.0000,
                    0.0000, 0.0000, 0.5189, 0.2862, 0.0000, 0.1509, 0.0000, 0.0000, 0.0000,
                    0.0000, 0.0000],
                    [0.0000, 0.0000, 0.0000, 0.0000, 0.2394, 0.0326, 0.4283, 0.2156, 0.0000,
                    0.0000, 0.0000, 0.6659, 0.1813, 0.0000, 0.1295, 0.0000, 0.0000, 0.0000,
                    0.1829, 0.0000],
                    [0.0000, 0.1029, 0.0000, 0.0000, 0.0000, 0.3352, 0.0000, 0.0000, 0.0000,
                    0.0000, 0.0912, 0.5666, 0.0000, 0.0000, 0.0430, 0.0000, 0.0000, 0.1878,
```

nn.Sequential

nn.Sequential is an ordered container of modules. The data is passed through all the modules in the same order as defined. You can use sequential containers to put together a quick network like `seq_modules`.

```
seq_modules = nn.Sequential(
    flatten,
    layer1,
    nn.ReLU(),
    nn.Linear(20, 10)
)
input_image = torch.rand(3,28,28)
logits = seq_modules(input_image)
```

nn.Softmax

The last linear layer of the neural network returns *logits* - raw values in [-inf, inf] - which are passed to the **nn.Softmax** module. The logits are scaled to values [0, 1] representing the model’s predicted probabilities for each class. `dim` parameter indicates the dimension along which the values must sum to 1.

```
softmax = nn.Softmax(dim=1)
pred_probab = softmax(logits)
```

Model Parameters

Many layers inside a neural network are *parameterized*, i.e. have associated weights and biases that are optimized during training. Subclassing `nn.Module` automatically tracks all fields defined inside your model object, and makes all parameters accessible using your model’s `parameters()` or `named_parameters()` methods.

In this example, we iterate over each parameter, and print its size and a preview of its values.

```
print(f"Model structure: {model}\n\n")

for name, param in model.named_parameters():
    print(f"Layer: {name} | Size: {param.size()} | Values : {param[:2]} \n")
```

Out:



```
Model structure: NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)

Layer: linear_relu_stack.0.weight | Size: torch.Size([512, 784]) | Values : tensor([[ 0.0215, -0.0305,  0.0335, ..., -0.0046, -0.0254,
 0.0260],
      [-0.0074,  0.0131, -0.0024, ...,  0.0203, -0.0169, -0.0319]],
      device='cuda:0', grad_fn=<SliceBackward0>)

Layer: linear_relu_stack.0.bias | Size: torch.Size([512]) | Values : tensor([ 0.0183, -0.0205], device='cuda:0', grad_fn=
<SliceBackward0>)
```

Further Reading

- [torch.nn API](#)

Total running time of the script: (0 minutes 0.108 seconds)

[< Previous](#)

[Next >](#)

Rate this Tutorial ☆☆☆☆☆

Docs

Access comprehensive developer documentation for PyTorch

[View Docs](#)

Tutorials

Get in-depth tutorials for beginners and advanced developers

[View Tutorials](#)

Resources

Find development resources and get your questions answered

[View Resources](#)

PyTorch

Get Started

Features

Ecosystem

Blog

Contributing

Resources

Tutorials

Docs

Discuss

Github Issues

Brand Guidelines

Stay Connected